

Course TÖL401G: Stýrikerfi / Operating Systems 13. I/O Systems

Mainly based on slides and figures
copyright Silberschatz, Galvin and Gagne

Chapter Objectives

- Explore the structure of an operating system's I/O subsystem.
- Discuss the principles and complexities of I/O hardware.
- Explain the performance aspects of I/O hardware and software.

Contents

1. Overview I/O Systems
2. I/O Hardware
3. *Application I/O Interface*
4. *Kernel I/O Subsystem*
5. *Transforming I/O Requests to Hardware Operations*
6. *Performance*
7. Summary

13.1 Overview

Input/Output Systems

- **Magnitude of I/O devices** (storage, transmission, human-interface):
 - Printer
 - Scanner
 - Keyboard
 - Mouse
 - Game controller
 - Digitiser touch screen/tablet
 - Hard disk
 - Flash memory SSD disk
 - CD/DVD/Blu-ray Reader & Writer
 - Tape drive
 - Graphic card
 - Sound card
 - GPS
 - Accelerometer
 - Wired & Wireless network
 - ...
- **Magnitude of hardware interfaces:**
 - USB
 - (Serial) ATA
 - NVMe Express (NVMe)
 - Serial-attached SCSI (SAS)
 - Thunderbolt,
 - PCI Express (PCIe),
 - Legacy interfaces (parallel, serial),
 - Ethernet, ...

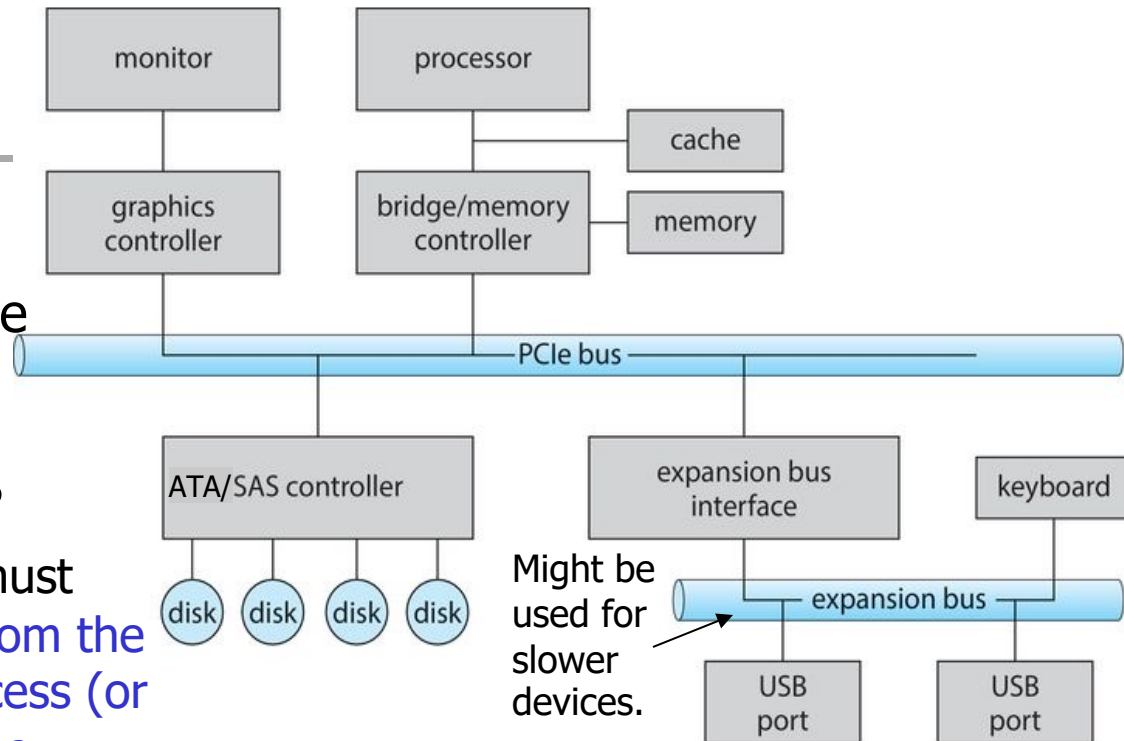
I/O and the Operating System

- Responsibility of the OS concerning I/O devices:
 - Manage I/O devices (using **device drivers**):
 - Provide processes/programmer a simple, but controlled access.
 - In multiprogramming systems:
 - Provide a fair access to I/O device used by multiple processes.
 - Minimise I/O performance bottleneck using suitable scheduling.
 - I/O devices are much much slower than CPU and main memory.
 - Avoid that I/O-intensive processes slow down CPU-intensive processes!
 - Cope with magnitude of different devices and interfaces:
 - Application developer should not have to deal with magnitude.
 - ⇒ Provide a standard interface for unified access to devices with comparable characteristics.

13.2 I/O Hardware

- How do device drivers (=software) communicate with the actual I/O devices or their device controllers (=hardware)?

- OS and I/O hardware must support **sending data from the address space of a process (or of the OS performing the respective I/O system call) to an I/O device and vice-versa.**
- Devices are either directly connected to the CPU via the data and address bus, or via some peripheral bus, e.g. the PCIe bus (having multiple PCIe "lanes").



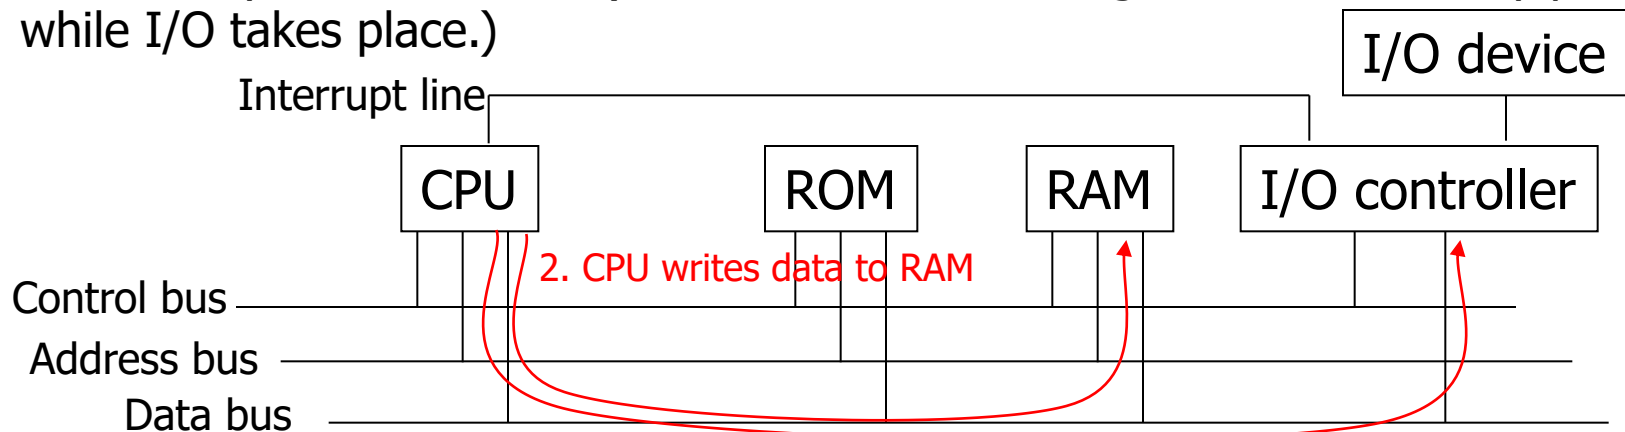
- Note: Nowadays, many controllers are not plugged into a PCIe slot, but integrated into the bridge/memory controller (it's "southbridge" part).
- In fact, since 2011, the bridge/memory controller (former "northbridge") is part of the CPU chip itself, and the device controllers (incl. integrated graphics) are located in the Platform Controller Hub (PCH) chip.
- Starting from 2014 also the PCH is sometimes part of the CPU chip itself.

I/O Hardware

- Possible ways for connecting an I/O controller to a CPU or main memory (more on next slides):
 - Programmed I/O (PIO),
 - Memory-mapped I/O,
 - Direct Memory Access (DMA).
- Knowing the different characteristics of these hardware connections may help minimising the I/O performance bottleneck.
 - Optimal: While I/O device is working (corresponding process is blocked), CPU can be assigned to other (ready) process.

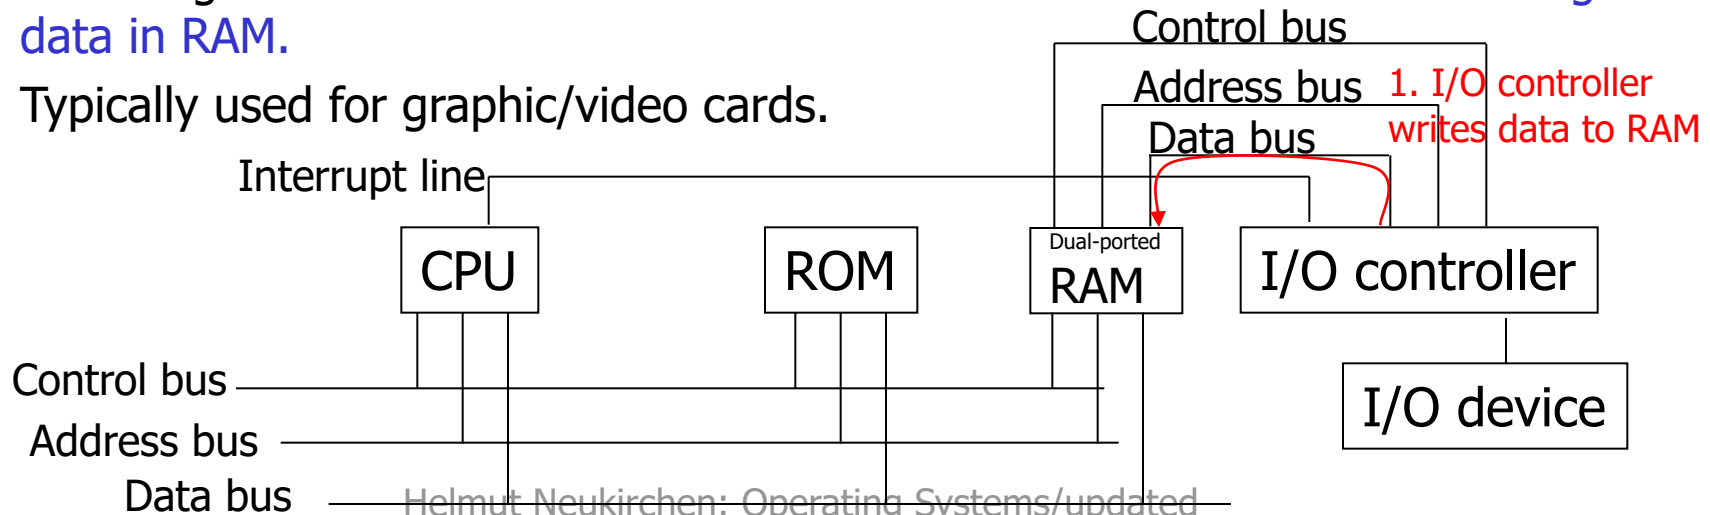
Programmed I/O (PIO)

- CPU is connected to a data register ("I/O port") of the controller, i.e. device driver or CPU respectively have to read/write data byte-wise from/to controller.
 - **Polling**: Device driver (=CPU) has to read periodically status register of controller whether data needs to be transferred from/to controller. (Disadvantage: **busy waiting**)
 - **Interrupt**: Controller raises interrupt if data must be transferred. Interrupt handler dispatches to device driver that transfers actual data from/to controller. (Disadvantage: interrupt handling is slow due to involved **context switch**.)
- Disadvantage of PIO: not suitable for multiprocessing, as byte-wise transfer of data is very **CPU intensive**! (I.e. CPU cannot be assigned to some ready process while I/O takes place.)



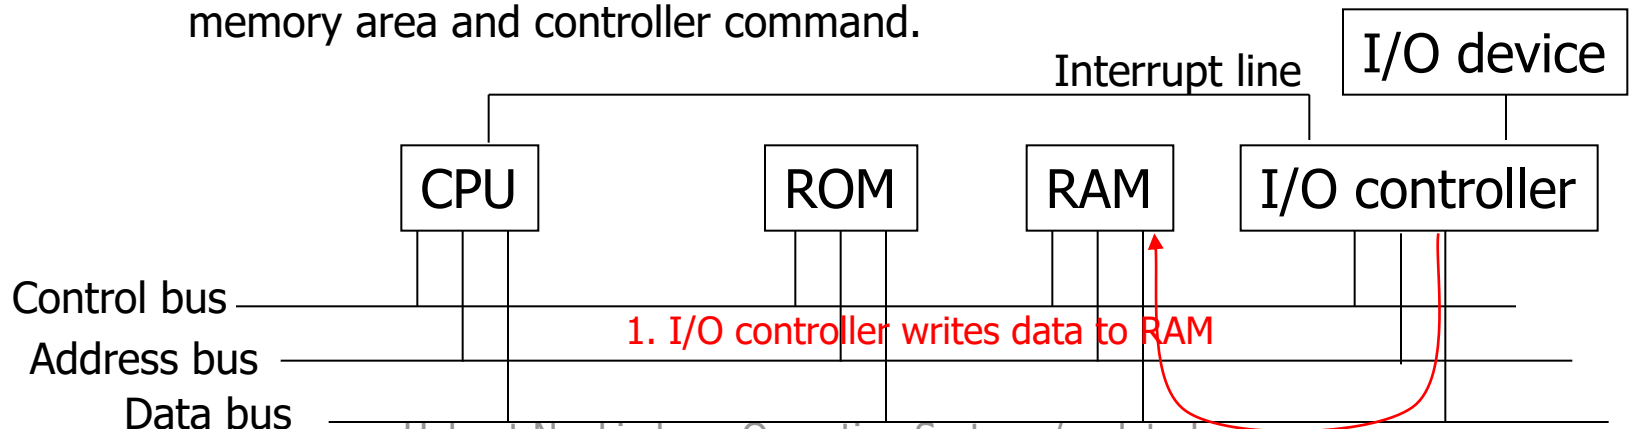
Memory-Mapped I/O

- Controller (e.g. graphic card) has its own RAM memory that is mapped into the physical address space of the CPU.
- Disadvantage: RAM memory that can be used by two parties at the same time (i.e. by controller & CPU), is either **more expensive** (dual-ported RAM) **or slower** (shared Memory: I/O controller steals bus memory cycles from CPU to access ordinary (=single ported) RAM – however no problem if CPU cache contains all data needed by CPU) than ordinary RAM.
- Advantage: Less CPU intensive. **CPU can work while controller is accessing data in RAM.**
- Typically used for graphic/video cards.



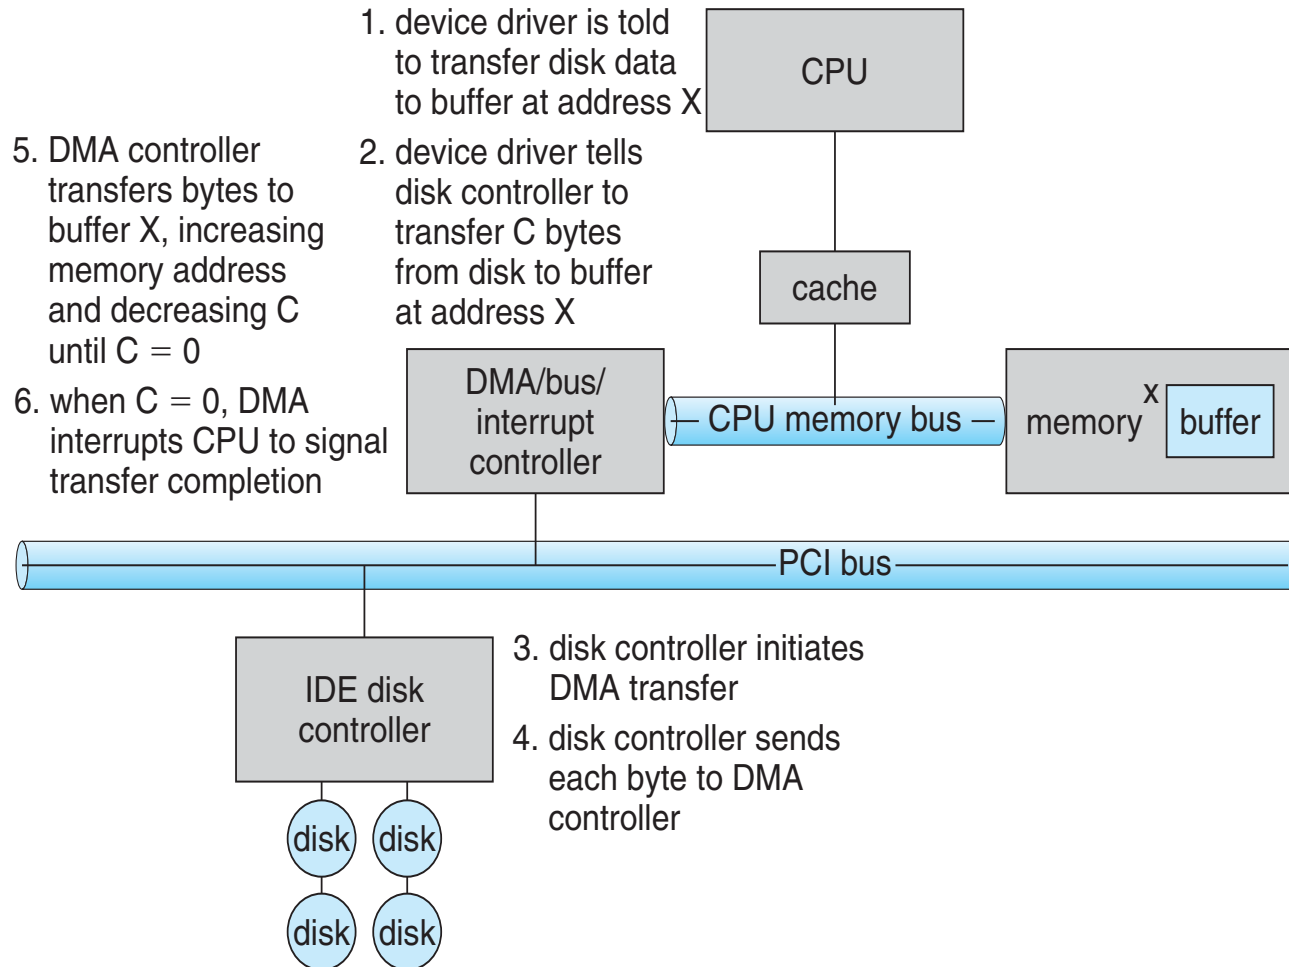
Direct Memory Access (DMA) (1)

- Controller (e.g. disk controller) is able to write from/to arbitrary addresses of the main memory without intervention of the CPU.
 - CPU just configures DMA controller once (using PIO) for each transfer:
 1. command to be executed by controller (e.g. sector to be read from disk & number of consecutive sectors to be read),
 2. start of memory address to read from/write to via DMA.
 - DMA controller raises interrupt when DMA transfer is finished.
 - CPU can then, e.g., issue another command.
 - Each DMA transfer comprises exactly one consecutive memory area and controller command.



Direct Memory Access (DMA) (2)

Detailed Steps



Direct Memory Access (DMA) (3)

Discussion

- Advantage: Less CPU intensive. CPU can work while controller accesses data in RAM.
 - Ideal for virtual memory using paging: While DMA controller is loading pages, CPU may still execute those processes that are ready and in physical memory.
 - However, frames that are currently subject of DMA transfer must not be chosen as victim for replacement while DMA transfer takes place!
 - Still, DMA transfer slows down access of CPU to RAM (“memory cycle stealing” by DMA controller to access RAM in parallel to CPU just like shared-memory approach of memory-mapped I/O).
 - But again: no slow down if data needed by CPU is in CPU cache.
- All major device controllers that transfer huge amounts of data are nowadays DMA-enabled.
 - E.g. hard disk/SSD controller, USB controller.

13.7 Summary

- DMA attachment of mass storage devices allows CPU to compute in parallel to I/O.
- OS provides a unified interface for accessing I/O devices.
- Layered approach allows to exchange low-level device drivers while keeping higher-level device-independent drivers and file system implementations.
- Kernel I/O subsystem may help to reduce I/O bottleneck.