

HBV401G SOFTWARE DEVELOPMENT

12. Software Testing

Matthias Book
Spring 2022

**FACULTY OF INDUSTRIAL ENGINEERING, MECHANICAL
ENGINEERING AND COMPUTER SCIENCE**

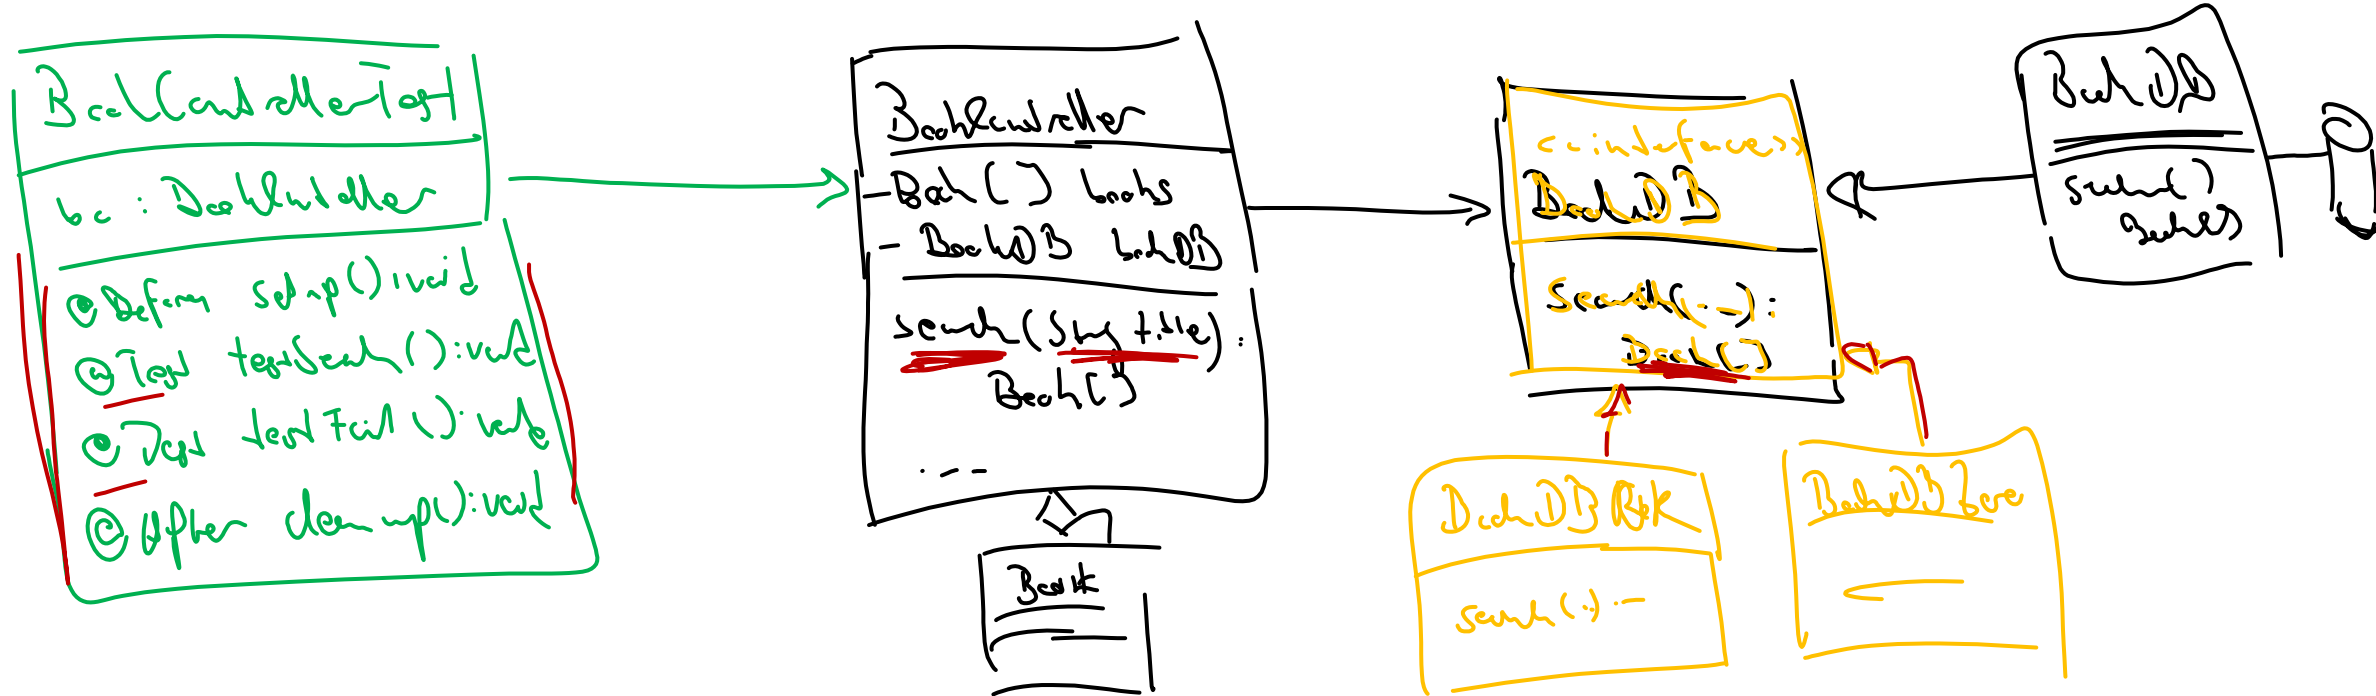
In-Class Quiz #11 Solution: Software Testing Concepts

- Indicate which concepts are described by the following definitions – Test Cases, Test Fixtures, Test Oracles, Test Subjects or Mock Objects?
- a) **Test Oracle:** An entity determining what a test subject's expected "correct" output for a particular input is supposed to be.
- b) **Test Fixture:** A piece of code exposing a test subject to a particular test case.
- c) **Mock Object:** A piece of code that simulates certain behavior of other, not yet implemented code that a test subject is relying on.
- d) **Test Subject:** A piece of code whose correctness / agreement with specifications is being tested.
- e) **Test Case:** A specific output/result that is expected to be produced by a test subject when exposed to a specific input.




Recap: Testing BookController.search()

Class Diagram



Recap: Testing BookController.search() Sequence Diagram



Testing Strategies



Testing a Method that Returns Nothing (e.g. a Constructor or `void` Return Type)

- A method returning nothing should still have some side effect (otherwise it's pointless!)
 - Test for the presence of the side effect!

- Example:

```
public class A {  
    private int x;  
    public void setX(int x) {  
        this.x = x;  
    }  
    public int getX() {  
        return x;  
    }  
}
```

```
@Test  
public void testSetX() {  
    a.setX(42);  
    assertEquals(42, a.getX());  
}
```

Testing for `setX`'s effect
on the attribute `x`

Should We Test Getters and Setters?

- If their implementation is trivial, you typically don't need to:

For a trivial implementation such as...

```
public class A {  
    private int x;  
    public void setX(int x) {  
        this.x = x;  
    }  
    public int getX() {  
        return x;  
    }  
}
```

```
@Test  
public void testGetSetX() {  
    a.setX(42);  
    assertEquals(42, a.getX());  
}
```

...is pretty much the same as...

```
@Test  
public void testGetSetX() {  
    assertEquals(42, 42);  
}
```

- But if they perform any input validation/output transformation, you likely want to.

Testing if Exceptions Are Raised

- Suppose a constructor is expected to raise an `IllegalArgumentException` if it receives an empty currency string:

```
@Test(expected=IllegalArgumentException.class)
public void testIllegalConstructorArgument() {
    Money undef500 = new Money(500, "");
}
```

Will fail if the exception is not raised in this situation.

- Note we need another test to ensure the exception is only thrown with reason!

```
@Test
public void testLegalConstructorArgument()
    throws IllegalArgumentException {
    Money isk500 = new Money(500, "ISK");
}
```

Will fail if the exception is raised in this situation.

Note: Only both tests together ensure correct exception behavior! One test alone could be satisfied by a method that always/never throws an exception, even though that would not be our desired behavior.

Testing for `null` References

- Should we test for proper treatment of `null` references in these methods?

```
public int convertTo(String toCurrency) {  
    return converter.convert(amount, currency, toCurrency);  
}
```

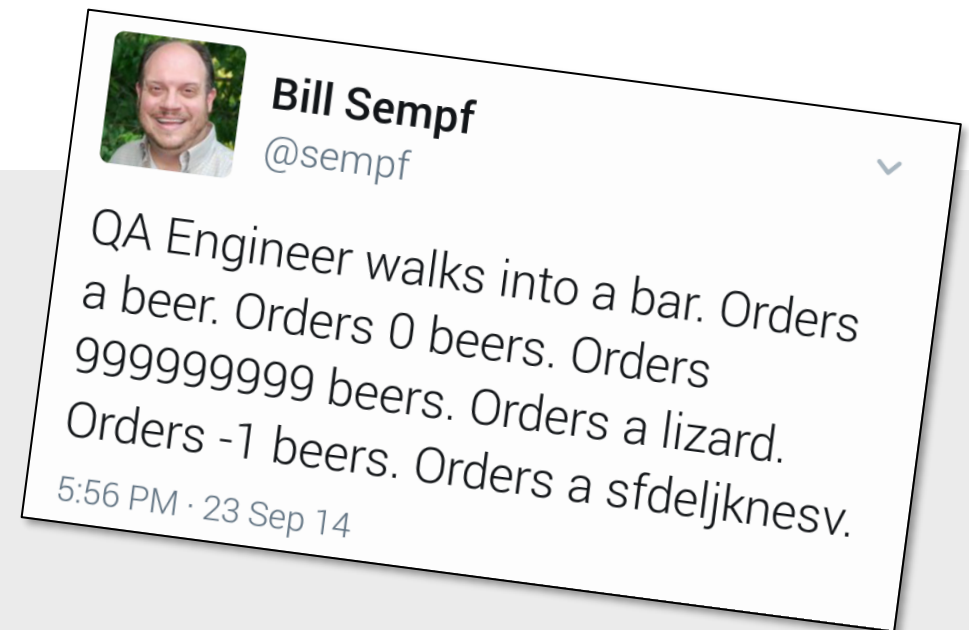
- No – making sure we have a reference to a `converter` object is the job of the constructor (or whoever provides the `converter`), and should be tested for there.
 - This method should be able to rely on that, and does not need to do any exception handling.

```
public Money add(Money summand) {  
    return new Money(amount + summand.getAmount(), currency);  
}
```

- No – if the `summand` is `null`, the JVM's default behavior of throwing a `NullPointerException` is the only sensible reaction.
 - As Java's default behavior, no explicit implementation or test is necessary at this point.
 - We may want to test that no `null` references can be produced at the `summand`'s origin though.

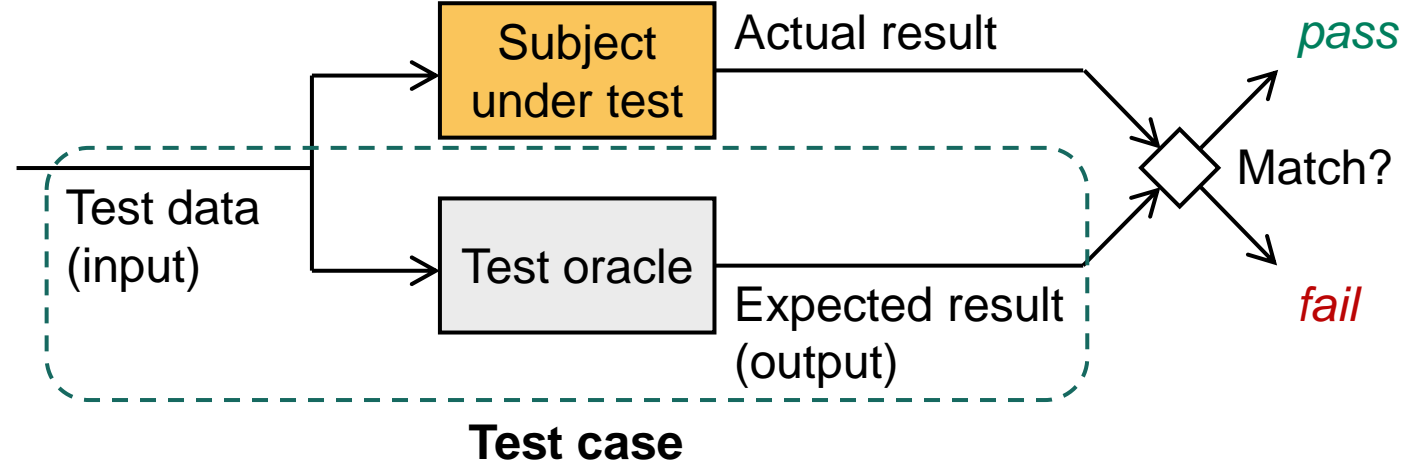
Test Coverage:

Equivalence Partitioning, Boundary Value Analysis



Choosing Appropriate Test Data

- The test subject is exposed to the test data, and the actual result is compared to the expected result.



- Question: How do we choose appropriate test data? – i.e:
 - Test data that covers all possible scenarios that the test subject could experience
 - Black-box view: Test cases designed to cover required behaviors
 - Test data that covers all possible ways in which the test subject can behave
 - White-box view: Test cases designed to cover implemented execution paths
- Problem: It is usually impractical/impossible to achieve 100% coverage.
 - What is the most economic subset to choose?
 - i.e. **how can we find the most defects with the least effort?**

e.g. equivalence partitioning, boundary analysis (*up next*)

various coverage criteria (*beyond scope of this course*)

Equivalence Partitioning

Note: These are not object-oriented classes, but mathematical sets of input values

- **Problem:** Typically, it is impractical/impossible to expose the test subject to all possible combinations of input values.
- **Idea:** Partition the set of all possible input values into “equivalence classes”, i.e. subsets for which the test subject’s behavior will be the same.
 - i.e. we want to identify subsets of input values for which the test subject will behave the same (e.g. a given test will pass for any member of the subset, or fail for any member).
 - Once we have identified these equivalence classes *for a particular test*, we only need to test one representative of each equivalence class instead of the whole set of input values.
- **Testing procedure**
 1. Determine conditions imposed on the input data (should be found in specification docs)
 2. Partition overall input value range into equivalence classes
 3. Pick a representative from each equivalence class to serve as test input

- For a given test, we want to find [combinations of] input values for which the test yields the same result (i.e. passes or fails).
- Various strategies can be used, depending on the kind of input value range:
- **Ordered input value range**
 - The range of acceptable values forms a “valid input” class
 - The ranges of values below and above the acceptable values form “invalid input” classes
 - Example:
 - Complete value range: $i \in \mathbb{Z}$, acceptable values: $1 \leq i \leq 10$
 - “valid input” equivalence class: $1 \leq i \leq 10$
 - “invalid input” equivalence classes:
 - $i < 1$
 - $i > 10$

■ Collections of objects

- Example: A list may contain between 1 and 25 elements
- “valid input” equivalence class: all lists containing between 1 and 25 elements
- “invalid input” equivalence classes:
 - an empty list
 - all lists containing over 25 elements
 - a `null` reference

■ Distinction of individual input values

- Example: Same behavior for inputs ‘a’, ‘b’ and ‘c’, but different behavior for ‘x’ and ‘y’
- “valid input” equivalence classes:
 - {‘a’, ‘b’, ‘c’}
 - {‘x’, ‘y’}
- “invalid input” equivalence class: all other inputs

- Pick / construct a combination of test input data that covers *as many “**valid input**” equivalence classes as possible*
 - (i.e. the [components of the] test case’s input data should be representative[s] of as many “valid input” equivalence classes as possible)
 - Repeat this as long as there are still “valid input” equivalence classes from which no representative has been used yet

- Pick / construct a combination of test input data that covers *one “**invalid input**” equivalence class*
 - Repeat this as long as there are still “invalid input” equivalence classes from which no representative has been used yet

Example: Equivalence Partitioning

Steps 1 & 2: Construct Equivalence Classes

- We want to test the method `String read(String filename, int n)` that is supposed to return the first `n` lines of the file with the given `filename`.
- **Step 1: Determine conditions imposed on input data:** According to specs,
 - the filename should consist of 1 to 6 characters (numbers or letters)
 - the first character of the filename should be a letter
 - the number of requested lines should be greater than 0 and less than 1000
- **Step 2: Partition equivalence classes**

Input condition	“Valid input” equiv. classes	“Invalid input ” equiv. classes
Filename length	1. 1 to 6 characters	a) Empty filename b) More than 6 characters
Filename characters	2. Numbers or letters	a) Other characters
1 st filename character	3. Letter	a) No letter
Number of lines	4. > 0 and < 1000	a) ≤ 0 b) ≥ 1000

Example: Equivalence Partitioning

Step 3: Construct Test Cases

Test case covering all “valid input” equivalence classes:

- `read("abc1", 42);` (1)–(4)

Test cases covering individual “invalid input” equivalence classes:

- `read("", 42);` (1a)
- `read("abc1234", 42);` (1b)
- `read("a+", 42);` (2a)
- `read("21a", 42);` (3a)
- `read("abc1", -21);` (4a)
- `read("abc1", 2100);` (4b)

Input condition	“Valid input” equiv. classes	“Invalid input ” equiv. classes
Filename length	1. 1 to 6 characters	a) Empty filename b) More than 6 characters
Filename characters	2. Numbers or letters	a) Other characters
1 st filename character	3. Letter	a) No letter
Number of lines	4. > 0 and < 1000	a) ≤ 0 b) ≥ 1000

How Representative are the Representatives?

- Mathematically, it seems sufficient to pick just one representative from each equivalence class:
 - If our code should accept $1 \leq x \leq 100$, isn't it sufficient to test just e.g. -50, 50, 150?
 - If our implementation is correct, these values are indeed representative of their eqv. classes
- However, if our implementation is defective, it might actually imply different equivalence classes!
 - If a defective implementation rejects the input 42 for some reason, testing with the input values -50, 50 and 150 does not detect this discontinuity!
- Equivalence partitioning cannot guarantee the absence of defects.
 - But (testing economics): Putting such a discontinuity into the code by accident is rare, so we typically can have reasonable confidence in the result.

- Still, we should be aware of parts of our code that could introduce such discontinuities, or that imply different equivalence classes than we intended, and test if we treated them correctly.
- Such defects can hide anywhere, but are most easily introduced inadvertently at the boundaries of our intended equivalence classes. Typical examples:
 - Loop counters
 - Array boundaries
 - Empty strings
 - Empty collections
 - Last days of a month
 - Overflows
 - `null` pointers
- Testing economics: Usually worth testing if boundaries are treated correctly!
 - So for the requirement $1 \leq x \leq 100$, we would e.g. test -50, 0, 1, 2, 50, 99, 100, 101, 150

Quiz #12: Equivalence Classes & Boundary Values

- We want to test the method `String read(String filename, int n)`, where
 - the filename should consist of 1 to 6 characters (numbers or letters)
 - the number of requested lines should be greater than 0 and less than 1000
- Indicate which statements are well-constructed test cases for the “valid input” or “invalid input” equivalence classes, and which are poorly constructed test cases:
 - a) `String s = read("file", -1);`
 - b) `String s = read("file", 0);`
 - c) `String s = read("file", 1);`
 - d) `String s = read("filename", 999);`
 - e) `String s = read("filename", 1000);`
 - f) `String s = read("filename", 1001);`




Break

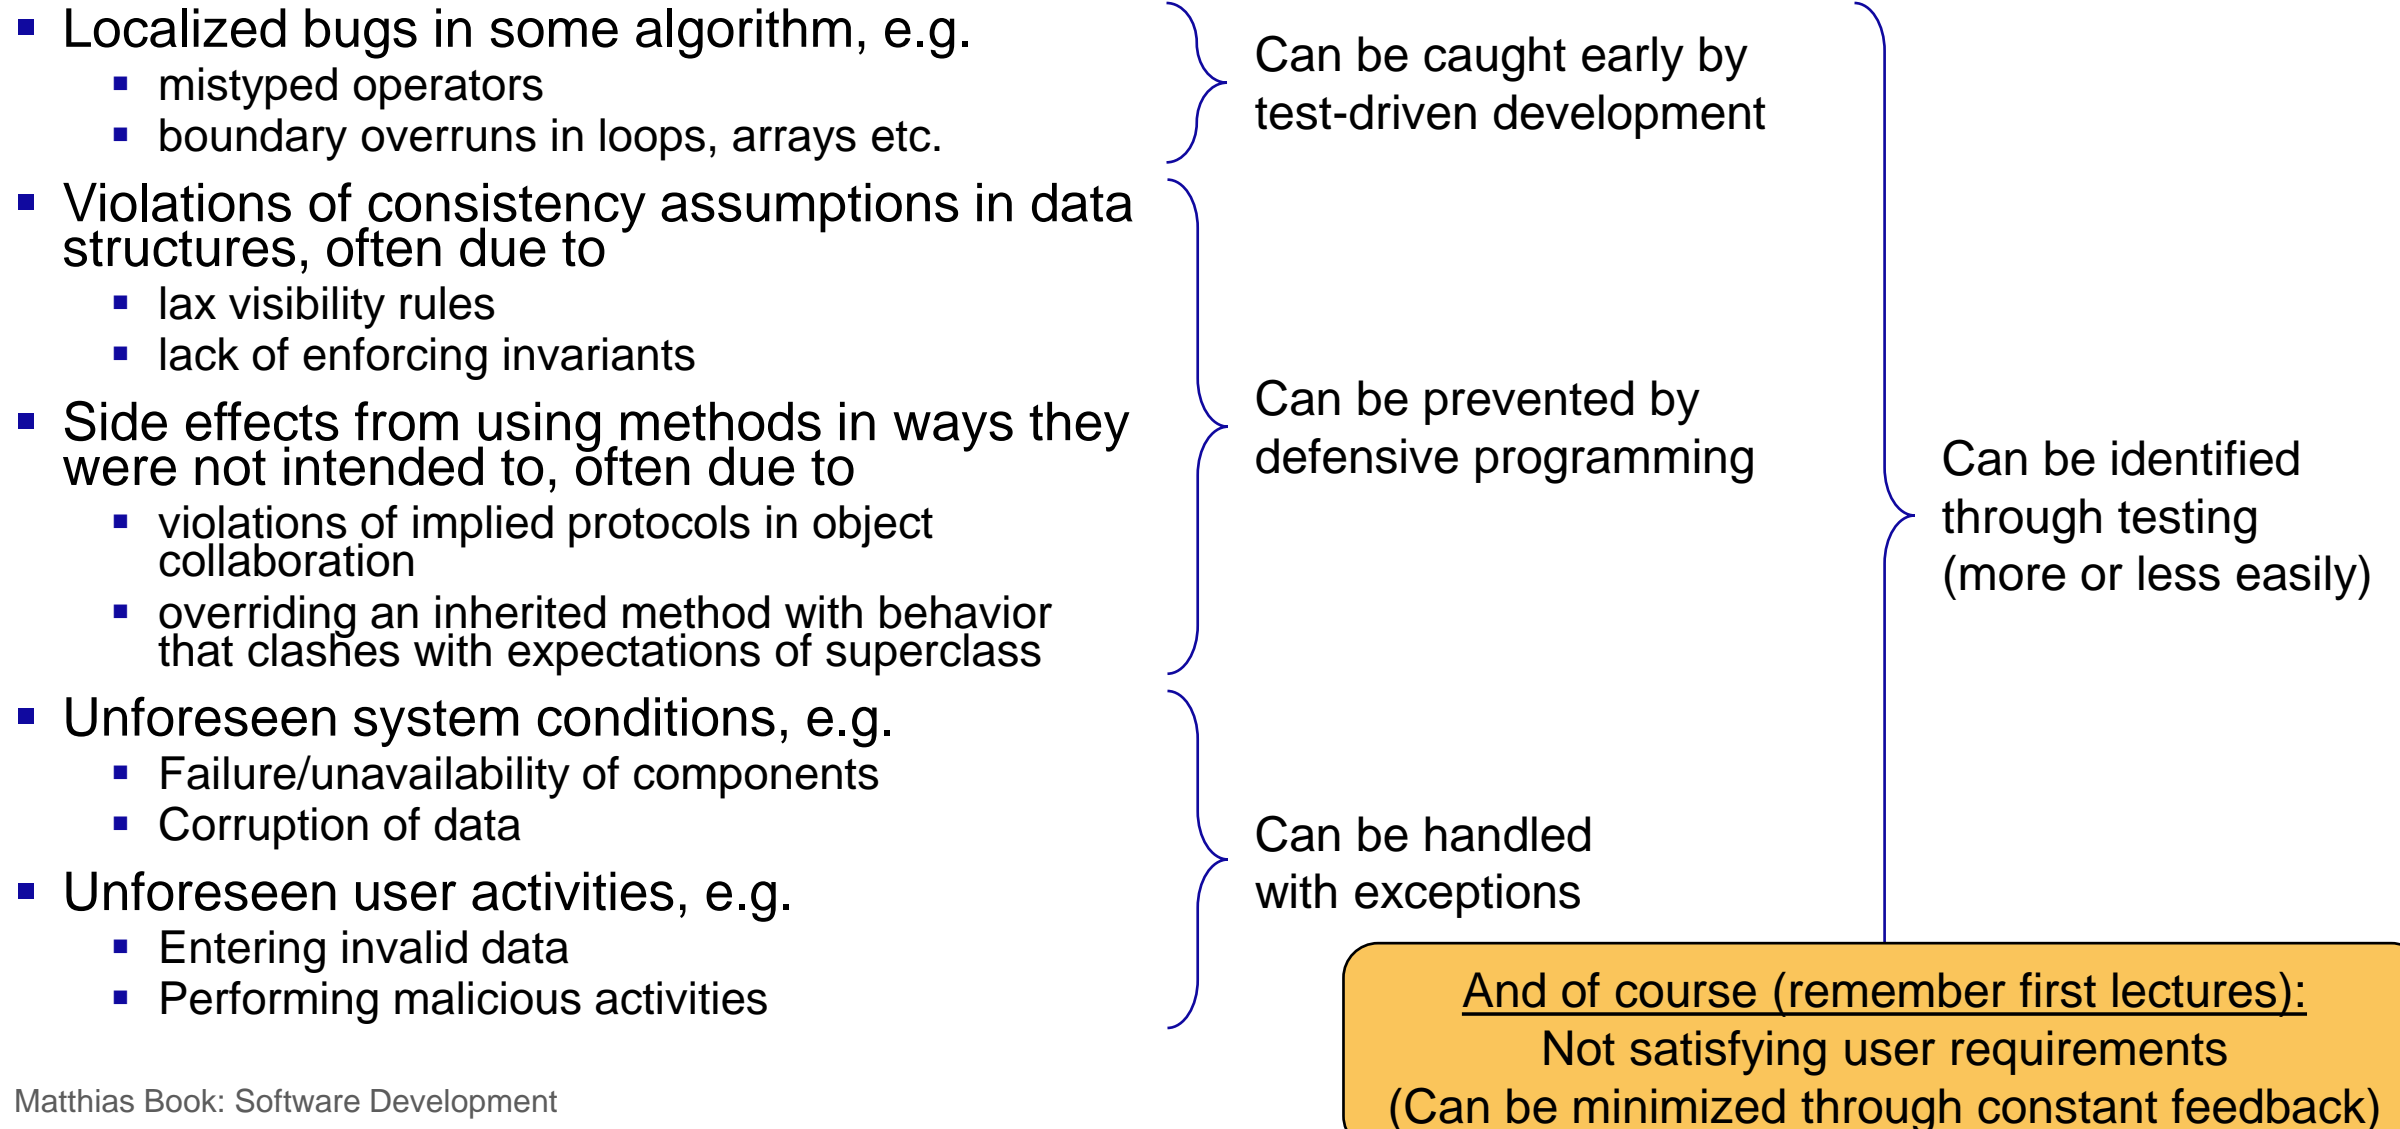


Defensive Programming

see also: J. Bloch, Effective Java



Why Programs Fail – And How to Prevent It



- **An object is responsible itself for always being in a valid state.**
 - i.e. its methods must guarantee that its attributes can't contain invalid [combinations of] data
- Use constructors to initialize attributes validly right from the creation of an object.
- Use setters to ensure that attributes are set within valid range.
- Make sure that the class' other methods maintain attributes within valid range.
- Check if incoming parameters are valid (within reason – “know who to trust”)
- Throw an exception if an operation would put your object in an invalid state.
- Use the **final** modifier on any attribute, parameter or variable that you don't intend to change after its initialization.
 - This is not just for constants, but can actually be applied to many parameters and variables.
 - Note that for reference types, **final** just means the object reference in the variable can't be changed anymore – the attributes within the object can still be changed.

Enforcing Consistency Through Encapsulation

- **Prefer working with objects over working with primitive data types.**
- Don't rebuild the reference (pointer) mechanism.
 - Don't give objects unique IDs and refer to them by ID – use the object references instead.
 - When dealing with objects coming from databases, you might still need the primary key to retrieve them – but consider whether there's a natural way to keep the object in memory between uses.
- Consider if it is more natural...
 - ...to pull lots of primitive data from object A and feed it to object B for processing
 - ...or to give object B a reference to object A and let it ask for the data it needs itself
- Within an object, store data in a way that balances...
 - ...the real-world structure of the application domain
 - This will keep your structure domain-oriented, technology-independent, and more maintainable
 - ...and the technical structures that are most efficient for working with the data
 - This will make data access more straightforward – but don't get too hung up on optimization!
- Use getters to transform data from your attributes as needed by other classes.
- Return empty collections instead of **null** references when a collection is expected.

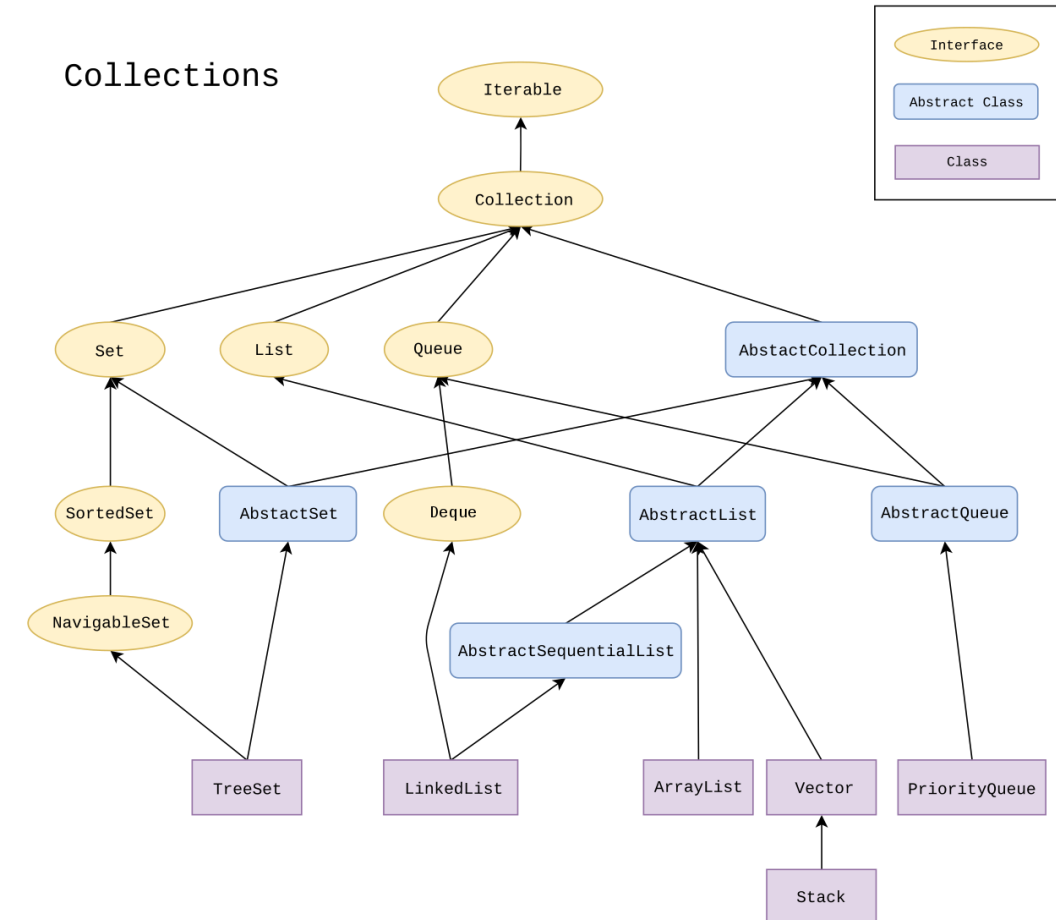
- **“Use the scope, Luke.” (Put any variable in the narrowest possible scope.)**
- Method variables should be used for any data that is only relevant within a method.
 - They should be declared in the block where they are required – this way they are destroyed after the block and can’t live on to create side effects
 - e.g. don’t declare a loop’s count variable at the beginning of the method, but in the loop head
- Instance variables (attributes) should be used only for data that is shared among the methods of one class, and that constitutes the state of the object.
 - Use **private** visibility to encapsulate attribute values within the object.
- Class variables (**static** attributes) are intended only for data that is shared among all instances of a class, or that shall be globally visible in the program.
 - This is very rarely necessary! Avoid unless you have a good reason.
 - Typically used e.g. for singletons or publicly visible constants
 - If you do this only because you don’t know how else to access some data from somewhere, rethink your design model!

- **Avoid any protocol beyond the method signature.**
- Other classes should normally not have to execute a particular set of your methods in a particular order to accomplish certain functionality.
 - Consider if you can do part of this automatically (e.g. by calling private methods).
- **Don't create a control language.**
 - Don't let methods react significantly differently to various input strings or input constellations.
 - Use overloaded methods instead; possibly rely on polymorphism.
 - Avoid signaling special conditions with special input/output parameters tucked into corners of the regular value range, e.g.
 - Special "optional" input parameters – consider overloaded methods, model "mode" explicitly
 - Special output values to signal error conditions – throw exceptions instead
- **Your method signatures should be...**
 - generic enough to give you the freedom of changing your internal implementation
 - close enough to your internal implementation to avoid awkward conversions
 - naturally representing the concepts of your application domain
- If you need to implement the same operation in similar ways, use an **interface** to describe the behavior, and implement the variants in several classes derived from it

- **Use methods in the context they are intended for.**
- Use visibility modifiers to make methods accessible only to intended audience.
- When overriding a method, ensure the new implementation satisfies all expectations the original implementation satisfied (Liskov Substitution Principle), i.e. behaves as expected...
 - when called from other methods that were inherited from the same superclass
 - when called by other classes who assume that this is an instance of the superclass
 - when calling other methods that the superclass' original method used to call
- Use **final** modifier to prevent a method from being overridden in subclasses.
 - Such methods can still be inherited, but not specialized anymore in subclasses.
- Consider using delegation / composition instead of inheritance to avoid some of the above access issues
 - Especially when the classes don't really have an "is-a" relationship, but just share code
 - Use inheritance when you want to make use of polymorphism

Example: Java Collections Framework

- Check out the documentation and API specification of the Java Collections Framework for good examples on how to encapsulate functionality, and how to design and document method signatures:
 - <https://docs.oracle.com/javase/8/docs/technotes/guides/collections/overview.html>
- The Collections Framework consists of
 - Collection interfaces
 - Collection implementations
 - General-purpose, wrapper, adapter, convenience, legacy, special-purpose, concurrent, abstract
 - Algorithms
 - Infrastructure
 - Utilities



Cluster Assignment #5: Final Product

(Presentations & Submissions for Clusters with a T Team)

- On **Wed 20 Apr**, present **your cluster's integrated product** to your classmates
 1. A demonstration of your cluster's integrated product (live, ~5 min)
 2. An overview of your cluster's system architecture (slides, ~5 min)
 3. A retrospective on your cluster's project work (slides, ~5 min)
 4. Q&A (~5 min)

- By **Sun 24 Apr**, submit in Canvas:
 1. A ZIP archive with the code of your cluster's integrated product
 2. A PDF document with the slides of your cluster's presentation

- Note the presentation comes before the submission for this assignment!
 - Don't fix/optimize your product after the presentation; your presented state counts!
 - Simply submit what you presented

Cluster Assignment #5: Final Product

(Presentations & Submissions for Clusters without a T Team)

- On **Wed 20 Apr**, present **your team's product** to your classmates
 1. A demonstration of your team's product (live, ~2 min)
 2. An overview of your team's system architecture (slides, ~2 min)
 3. A retrospective on your team's project work (slides, ~2 min)
 4. Joint Q&A with other teams in cluster (~2 min)

- By **Sun 24 Apr**, submit in Canvas:
 1. A ZIP archive with the code of your team's component
 2. A PDF document with the slides of your team's presentation

- Note the presentation comes before the submission for this assignment!
 - Don't fix/optimize your product after the presentation; your presented state counts!
 - Simply submit what you presented

■ Grading criteria

- **Demo:** Key features (searching, booking, reduced availability) are functional and accessible via the UI (75%)
- **Architecture:** Clear and competent description of the system's design (12.5%)
- **Process:** Clear and critical retrospective on the development process, discussion of lessons learned (12.5%)

■ Grading policy:

- Visible part of presentation → Assignment 5 grade for whole cluster/team
- Audible part of presentation
 - Given by one person → that person's Presentation Grade
 - Given by several people → contributes to Assignment 5 cluster/team grade

Cluster Assignment #5: Presentation Schedule

- Clusters present their products in the Zoom classroom at <https://eu01web.zoom.us/j/62847273071> on **Wed 20 Apr**
 - 15:00-15:20 Cluster 1
 - 15:25-15:45 Cluster 2
 - 15:50-16:10 Cluster 3
 - 16:15-16:35 Cluster 4
 - 16:40-17:00 Cluster 5
 - 17:05-17:25 Cluster 6
 - 17:30-17:50 Cluster 7
 - 17:55-18:15 Cluster 8
- All teams are encouraged to attend presentations of other clusters as well
 - to learn from other teams' experiences
 - to give other teams an audience

Cluster Assignment #5: Presentation Protocol

- Have your camera and microphone off by default
- Select “Hide non-video participants” in Zoom to see only the team “on stage”
- Turn your camera on (if available) during your own cluster’s/team’s presentation
- Speaker turns microphone on and shares screen to show your demo and slides
- Turn your microphone on during your team’s Q&A

- The final presentations will **not** be recorded.

Remaining Class Schedule

- **Wed 30 Mar** Assignment #4 (Software Test) presentations
- **Mon 4 Apr** Final lecture with information on final exam
- **Wed 6 Apr** Assignment #5 (Final Product) draft consultations
- **Mon 11 Apr** Spare lecture slot – anything you'd like to recap?
- **Wed 13 Apr** Easter break – no consultations
- **Mon 18 Apr** Easter break – no lecture
- **Wed 20 Apr** Assignment #5 (Final Product) presentations
- **Sun 24 Apr** Assignment #5 (Final Product) and peer assessment submission
- **Tue 26 Apr** Final exam

Thank you!

book@hi.is

